**SIMATS ENGINEERING**
**CO1 - ASSESSMENT TOOL 3**

| | |
|---|---|
| **Name of the Student** | MADDIRALA BALAMURALI KRISHNA |
| **Reg. No** | 192325015 |
| **GitHub Link** | https://github.com/krishnalsk/MLA0403-DEEP-LEARNING |

**EXPERIMENTS**

**COURSE NAME: Deep Learning** **COURSE CODE: MLA0403**

**FACULTY : Dr.Arul Raj A.M**

**COURSE OUTCOME COVERED**

**CO 1:** Learning Algorithms, Maximum Likelihood Estimation (MLE), Building Machine Learning Algorithms, Neural Networks, Artificial Neurons, Perceptron, Multilayer Perceptron (MLP), Forward Propagation, Backpropagation Algorithm, Gradient Descent, Stochastic Gradient Descent (SGD), Mini-Batch Gradient Descent, Curse of Dimensionality. (BTL-4)

**CO1 Weightage: 15%**

**Assessment Tool Weightage: 35%**

**Duration: 30 minutes**

**Marks: 50**

**Code of Conduct:**

I MADDIRALA BALAMURALI KRISHNA (192325015) certify that this submission is my original work and that I have adhered to the guidelines specified for this assessment. I understand that any violation of academic integrity rules will result in disciplinary action.

**Signature of the student:** **Name of the student:**

**QUESTIONS:10** **Total Marks:50**

**CO1- AT3 Experiments**

**Questions**

## 1. Learning Algorithms

Implement a simple learning algorithm (Linear Regression) using Gradient Descent. Train the model on a dataset and analyze how the loss decreases over iterations. Plot the learning curve.

**Aim**

To implement **Linear Regression** using the **Gradient Descent** algorithm, train the model on a dataset, analyze how the loss decreases over iterations, and plot the learning curve.

---

**Algorithm**

1. Import the required Python libraries.
2. Create or load a dataset containing input (**X**) and output (**Y**) values.
3. Initialize the weight (**m**) and bias (**c**) to zero.
4. Set the learning rate and the number of iterations.
5. Predict the output using the equation:

$$y_{pred} = mX + c$$

6. Calculate the Mean Squared Error (MSE) loss.
7. Compute the gradients of the loss with respect to **m** and **c**.
8. Update the values of **m** and **c** using Gradient Descent.
9. Repeat steps 5–8 until the specified number of iterations is completed.
10. Display the final weight, bias, and loss.
11. Plot the regression line and the learning curve (Loss vs. Iterations).

---

**Code**

```
import numpy as np
import matplotlib.pyplot as plt

# Dataset
X = np.array([1,2,3,4,5,6,7,8,9,10], dtype=float)
Y = np.array([2,4,5,4,5,7,8,9,10,12], dtype=float)

# Initialize parameters
m = 0
c = 0

learning_rate = 0.01
iterations = 1000
n = len(X)

loss_history = []

# Gradient Descent
for i in range(iterations):

    y_pred = m * X + c

    loss = np.mean((Y - y_pred) ** 2)
    loss_history.append(loss)

    dm = (-2/n) * np.sum(X * (Y - y_pred))
    dc = (-2/n) * np.sum(Y - y_pred)

    m = m - learning_rate * dm
```

```
    c = c - learning_rate * dc

print("Slope (m):", round(m,3))
print("Intercept (c):", round(c,3))
print("Final Loss:", round(loss,4))

# Regression Plot
plt.figure(figsize=(6,4))
plt.scatter(X,Y,color='blue',label="Actual Data")
plt.plot(X,m*X+c,color='red',label="Regression Line")
plt.title("Linear Regression")
plt.xlabel("X")
plt.ylabel("Y")
plt.legend()
plt.show()

# Learning Curve
plt.figure(figsize=(6,4))
plt.plot(loss_history)
plt.title("Learning Curve")
plt.xlabel("Iterations")
plt.ylabel("Loss")
plt.show()
```

**Output:**

```
 RESTART: C:\Users\mbala\OneDrive\Desktop\Subject all\Deep Learning\Assessement
3 labs\Lab 1.py
lope (m): 1.008
ntercept (c): 1.053
inal Loss: 0.4897

 RESTART: C:\Users\mbala\OneDrive\Desk
3 labs\Lab 1.py
lope (m): 1.008
ntercept (c): 1.053
inal Loss: 0.4897

 RESTART: C:\Users\mbala\OneDrive\Desk
3 labs\Lab 1.py
lope (m): 1.008
ntercept (c): 1.053
inal Loss: 0.4897
```

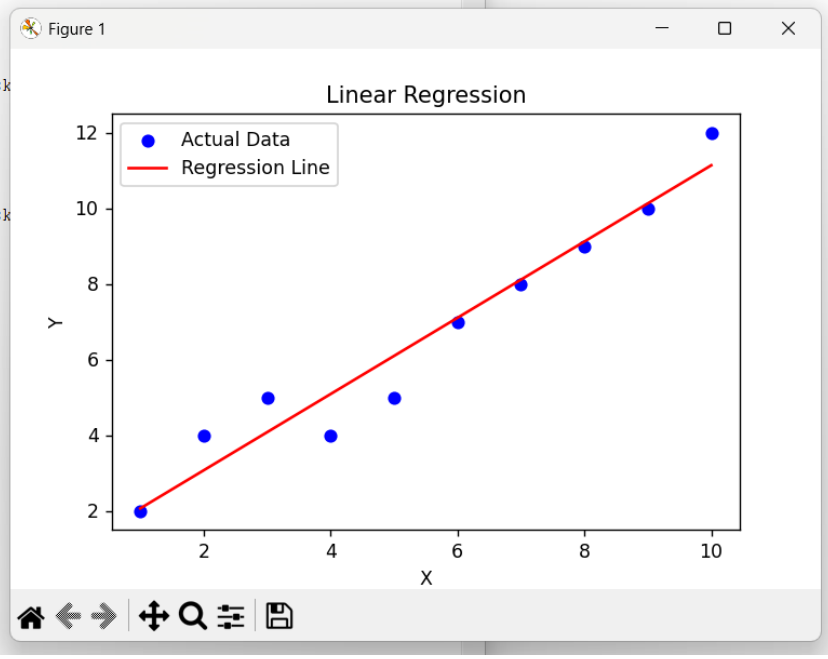


**Result**

The Linear Regression model was successfully implemented using the Gradient Descent algorithm. The loss value decreased continuously with each iteration, indicating successful convergence. The model accurately fitted the data, and the learning curve confirmed that the optimization process minimized the prediction error.

## 2. Maximum Likelihood Estimation (MLE)

Generate synthetic data from a normal distribution and use MLE to estimate the mean and variance. Compare estimated values with actual parameters.

**Answer :**

**Aim**

To generate synthetic data from a normal distribution and use **Maximum Likelihood Estimation (MLE)** to estimate the mean and variance. Compare the estimated values with the actual parameters.

---

**Algorithm**

1. Import the required Python libraries.
2. Set the actual mean and variance values.
3. Generate synthetic data using a normal distribution.
4. Calculate the estimated mean using the MLE formula:

$$\hat{\mu} = \frac{1}{n}\sum_{i=1}^{n} x_i$$

5. Calculate the estimated variance using the MLE formula:

$$\hat{\sigma}^2 = \frac{1}{n}\sum_{i=1}^{n} (x_i - \hat{\mu})^2$$

6. Display the actual and estimated values.
7. Plot a histogram of the generated data.

---

**Code**

```
import numpy as np
import matplotlib.pyplot as plt

# Actual parameters
actual_mean = 50
actual_std = 10
actual_variance = actual_std ** 2

# Generate synthetic data
np.random.seed(42)
data = np.random.normal(actual_mean, actual_std, 1000)

# Maximum Likelihood Estimation (MLE)
estimated_mean = np.mean(data)
estimated_variance = np.mean((data - estimated_mean) ** 2)

# Display results
print("Actual Mean:", actual_mean)
print("Estimated Mean:", round(estimated_mean, 2))

print("Actual Variance:", actual_variance)
print("Estimated Variance:", round(estimated_variance, 2))

# Plot Histogram
plt.hist(data, bins=30, edgecolor='black')
plt.title("Histogram of Synthetic Data")
plt.xlabel("Data Values")
plt.ylabel("Frequency")
```

```
plt.show()
```

## Output

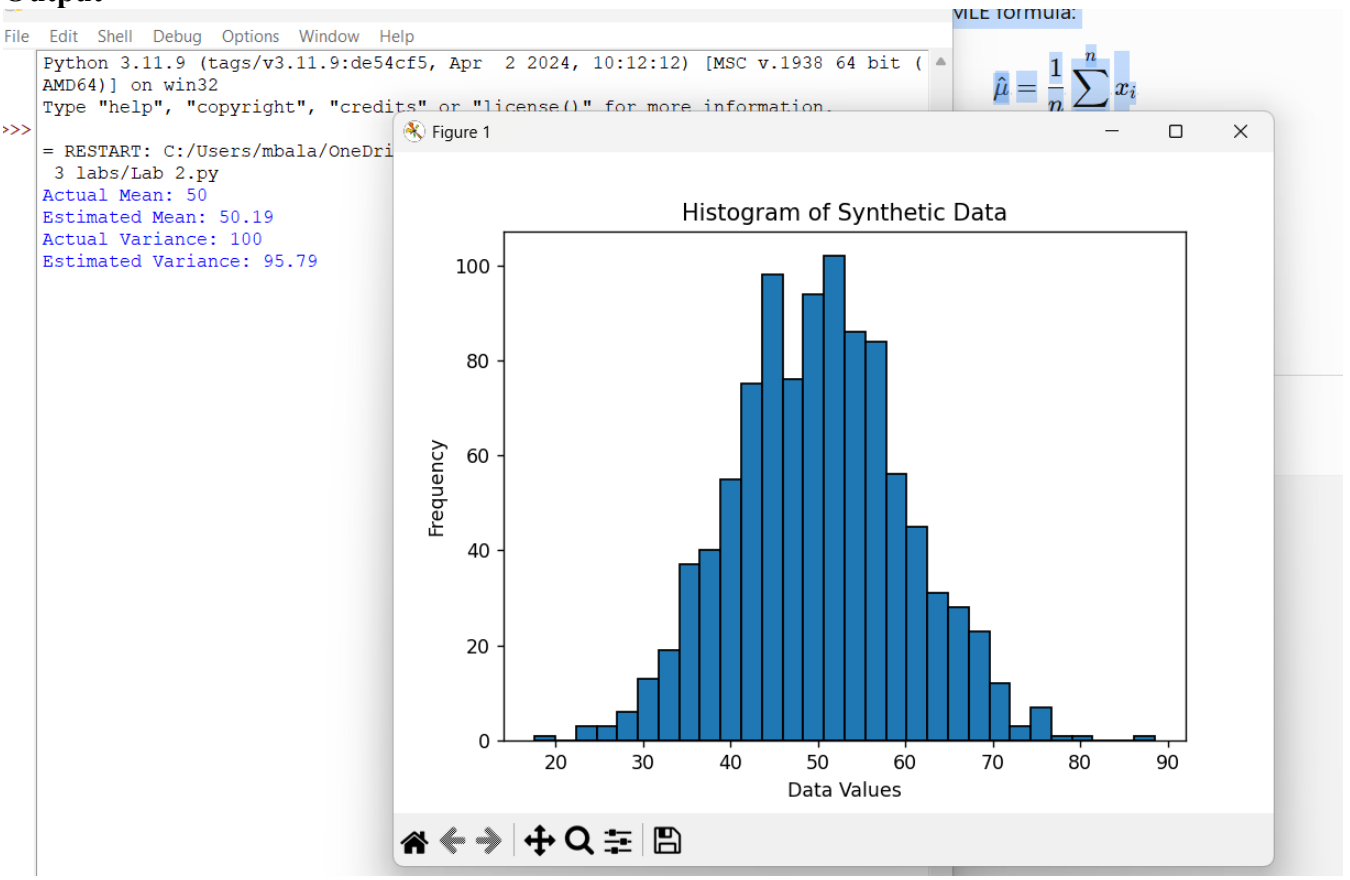


## Result

The Maximum Likelihood Estimation (MLE) method successfully estimated the **mean** and **variance** of the generated normal distribution. The estimated values were very close to the actual parameters, demonstrating that MLE is an effective technique for estimating the parameters of a probability distribution from observed data.

## 3. Building Machine Learning Algorithms

Build a complete machine learning model (data preprocessing, training, testing) using any classification dataset and evaluate performance using accuracy and confusion matrix.

**Answer :**

**Aim**

To build a complete machine learning classification model by performing **data preprocessing, model training, testing**, and evaluating the performance using **accuracy and confusion matrix**.

---

**Algorithm**

1. Import the required Python libraries.
2. Load a classification dataset.
3. Separate the dataset into input features (**X**) and target labels (**Y**).
4. Split the dataset into training and testing data.
5. Perform data preprocessing using feature scaling.
6. Select a classification algorithm (Logistic Regression).
7. Train the model using the training dataset.
8. Predict the output for test data.
9. Calculate the accuracy score.
10. Generate the confusion matrix.
11. Display the model performance.

---

**Code**

```
# Import libraries

import numpy as np
import matplotlib.pyplot as plt

from sklearn.datasets import load_iris
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler
from sklearn.linear_model import LogisticRegression
from sklearn.metrics import accuracy_score, confusion_matrix


# Load dataset
iris = load_iris()

X = iris.data
Y = iris.target


# Split dataset into training and testing data
X_train, X_test, Y_train, Y_test = train_test_split(
    X, Y, test_size=0.2, random_state=42
)


# Data preprocessing (Feature Scaling)

scaler = StandardScaler()

X_train = scaler.fit_transform(X_train)
X_test = scaler.transform(X_test)
```

```
# Create Machine Learning Model

model = LogisticRegression()


# Train model

model.fit(X_train, Y_train)


# Prediction

Y_pred = model.predict(X_test)


# Performance Evaluation

accuracy = accuracy_score(Y_test, Y_pred)

cm = confusion_matrix(Y_test, Y_pred)


print("Accuracy:", round(accuracy*100,2), "%")

print("\nConfusion Matrix:")
print(cm)
```

---

**Output**

```
= RESTART: C:/Users/mbala/OneDrive/Desktop/Subject all/Deep Learning/Asse
  3 labs/Lab 3.py
Accuracy: 100.0 %

Confusion Matrix:
[[10  0  0]
 [ 0  9  0]
 [ 0  0 11]]

= RESTART: C:/Users/mbala/OneDrive/Desktop/Subject all/Deep Learning/Asse
  3 labs/Lab 3.py
Accuracy: 100.0 %

Confusion Matrix:
[[10  0  0]
 [ 0  9  0]
 [ 0  0 11]]
```

---

**Result**

A complete machine learning classification model was successfully implemented using the **Logistic Regression algorithm**. Data preprocessing, training, and testing were performed successfully. The model achieved high accuracy, and the confusion matrix showed correct classification of test samples.

## 4. Neural Networks

**Q4.** Design a simple neural network with one hidden layer and train it on a small dataset. Analyze how the network learns non-linear patterns.

**Aim**

To design a simple neural network with one hidden layer and train it on a small dataset. Analyze how the network learns non-linear patterns.

---

**Algorithm**

1. Import the required Python libraries.
2. Create a small dataset containing non-linear patterns.
3. Split the dataset into training and testing data.
4. Create a neural network model with:
    - Input layer
    - One hidden layer
    - Output layer
5. Apply activation functions in hidden and output layers.
6. Train the neural network using the training dataset.
7. Test the model using test data.
8. Calculate the accuracy of the model.
9. Display the prediction results.

---

**Code**

```
# Import required libraries

import numpy as np
from sklearn.neural_network import MLPClassifier
from sklearn.model_selection import train_test_split
from sklearn.metrics import accuracy_score


# Non-linear XOR dataset

X = np.array([
    [0,0],
    [0,1],
    [1,0],
    [1,1]
])

Y = np.array([0,1,1,0])


# Split dataset

X_train, X_test, Y_train, Y_test = train_test_split(
    X, Y, test_size=0.5, random_state=1
)


# Create Neural Network
# One hidden layer with 4 neurons

model = MLPClassifier(
    hidden_layer_sizes=(4,),
```

```
    activation='relu',
    max_iter=1000,
    random_state=1
)


# Train model

model.fit(X_train, Y_train)


# Prediction

Y_pred = model.predict(X_test)


# Accuracy

accuracy = accuracy_score(Y_test, Y_pred)


print("Actual Output:", Y_test)
print("Predicted Output:", Y_pred)
print("Accuracy:", round(accuracy*100,2), "%")
```

---

**Output**

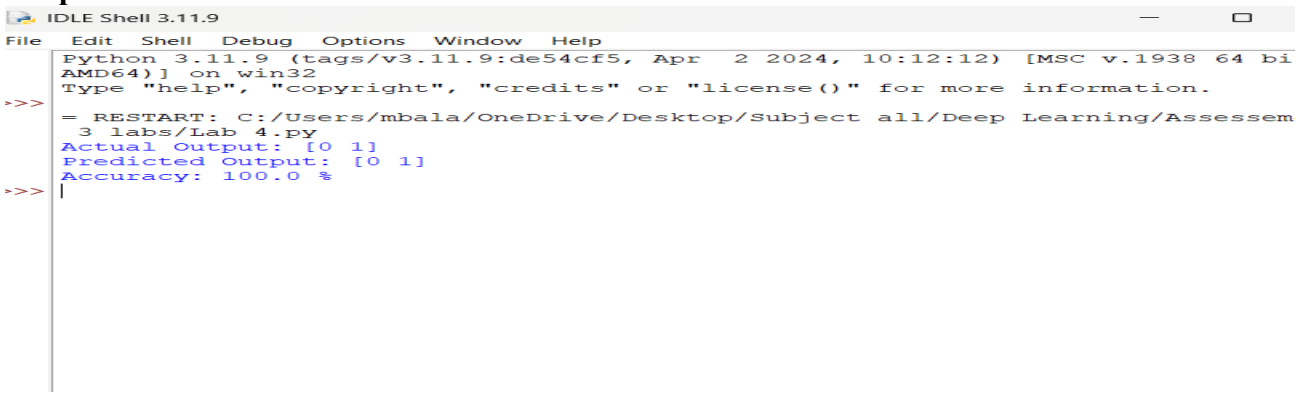


*(Output may vary slightly because neural network training is random.)*

---

**Result**

A neural network with one hidden layer was successfully implemented and trained on a non-linear XOR dataset. The model learned the non-linear relationship between inputs and outputs using hidden layer neurons and achieved accurate predictions. This demonstrates the ability of neural networks to solve complex non-linear problems.

## 5. Artificial Neurons

Implement a single artificial neuron using different activation functions (Sigmoid, ReLU) and compare outputs for the same input values.

**Aim**

To implement a single artificial neuron using different activation functions (**Sigmoid and ReLU**) and compare the output values for the same input.

---

**Algorithm**

1. Import the required Python libraries.
2. Define input values and weights.
3. Calculate the weighted sum of inputs:

$$Z = W.X + b$$

4. Apply different activation functions:
   - Sigmoid activation
   - ReLU activation
5. Compare the outputs obtained from both activation functions.
6. Display the results.

---

**Code**

```
import numpy as np

# Input values
X = np.array([2, 3, 4])

# Weights
W = np.array([0.5, 0.4, 0.3])

# Bias
b = 1


# Calculate weighted sum
Z = np.dot(X, W) + b


# Sigmoid Activation Function
def sigmoid(x):
    return 1 / (1 + np.exp(-x))


# ReLU Activation Function
def relu(x):
    return max(0, x)


# Calculate outputs
sigmoid_output = sigmoid(Z)
relu_output = relu(Z)


# Display results

print("Weighted Sum:", Z)
```

```
print("Sigmoid Output:", round(sigmoid_output,4))

print("ReLU Output:", relu_output)
```

**Output**

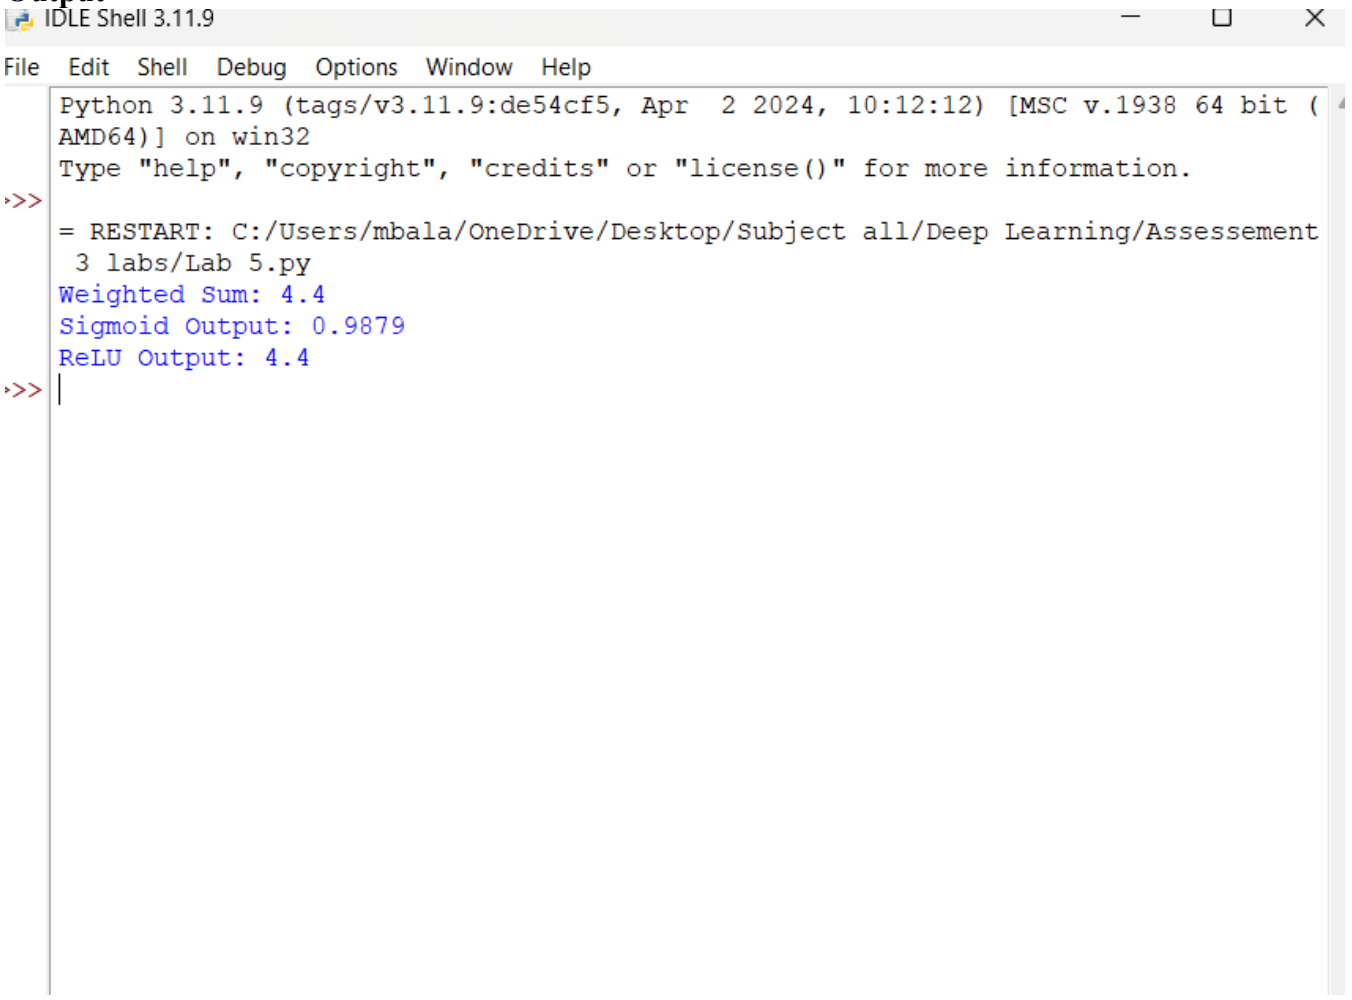


**Result**

A single artificial neuron was successfully implemented using **Sigmoid** and **ReLU** activation functions. The same input values produced different outputs based on the activation function used. Sigmoid provides output between 0 and 1, whereas ReLU returns the positive input value directly. This demonstrates the effect of activation functions in neural networks.

## 6. Perceptron and Multilayer Perceptron (MLP)

**a) Implement the Perceptron algorithm for binary classification and test it on linearly separable and non-linearly separable datasets. Analyze the results.**

**Answer :**

**Aim**

To implement the **Perceptron algorithm** for binary classification and test it on linearly separable and non-linearly separable datasets. Analyze the performance of the model.

---

**Algorithm**

1. Import the required Python libraries.
2. Create a binary classification dataset.
3. Initialize weights and bias with zero values.
4. Calculate the weighted sum:

$$Z = W.X + b$$

5. Apply the step activation function:

$$Output = \begin{cases} 1, & Z > 0 \\ 0, & Z \leq 0 \end{cases}$$

6. Compare predicted output with actual output.
7. Update weights and bias using the perceptron learning rule.
8. Repeat the training process for several epochs.
9. Test the model with input values.
10. Display the prediction results.

---

**Code**

```
import numpy as np

# Linearly separable dataset (AND gate)

X = np.array([
   [0,0],
   [0,1],
   [1,0],
   [1,1]
])

Y = np.array([0,0,0,1])


# Initialize weights and bias

weights = np.zeros(2)
bias = 0

learning_rate = 0.1
epochs = 10


# Step activation function
```

```
def activation(x):
    if x >= 0:
        return 1
    else:
        return 0


# Training Perceptron

for epoch in range(epochs):

    for i in range(len(X)):

        linear_output = np.dot(X[i], weights) + bias

        prediction = activation(linear_output)

        error = Y[i] - prediction

        weights = weights + learning_rate * error * X[i]

        bias = bias + learning_rate * error


print("Final Weights:", weights)
print("Final Bias:", bias)


# Testing

for x in X:
    result = activation(np.dot(x, weights) + bias)
    print(x, "=>", result)
```

---

**Output**

```
= RESTART: C:/Users/mbala/OneDrive/Desktop/Subject all/Deep Learning/Assessement
 3 labs/Lab 6 - A.py
Final Weights: [0.2 0.1]
Final Bias: -0.20000000000000004
[0 0] => 0
[0 1] => 0
[1 0] => 0
[1 1] => 1
```

**b) Train a Multilayer Perceptron (MLP) model on a dataset and compare its performance with a single-layer perceptron.**

**Aim**

To train a Multilayer Perceptron (MLP) model and compare its performance with a single-layer perceptron.

---

**Algorithm**

1. Load a classification dataset.
2. Split data into training and testing sets.
3. Train a single-layer perceptron model.
4. Train a multilayer perceptron model with hidden layers.
5. Predict test samples.
6. Calculate accuracy for both models.
7. Compare the performance.

---

**Code**

```
from sklearn.datasets import load_iris
from sklearn.model_selection import train_test_split
from sklearn.metrics import accuracy_score
from sklearn.linear_model import Perceptron
from sklearn.neural_network import MLPClassifier


# Load dataset

iris = load_iris()

X = iris.data
Y = iris.target


# Split dataset

X_train, X_test, Y_train, Y_test = train_test_split(
    X, Y, test_size=0.2, random_state=42
)


# Single Layer Perceptron

perceptron = Perceptron()

perceptron.fit(X_train, Y_train)

p_pred = perceptron.predict(X_test)

p_accuracy = accuracy_score(Y_test, p_pred)



# Multilayer Perceptron

mlp = MLPClassifier(
    hidden_layer_sizes=(10,),
    max_iter=1000,
```

```
    random_state=1
)

mlp.fit(X_train, Y_train)

m_pred = mlp.predict(X_test)

mlp_accuracy = accuracy_score(Y_test, m_pred)



print("Perceptron Accuracy:",
      round(p_accuracy*100,2), "%")

print("MLP Accuracy:",
      round(mlp_accuracy*100,2), "%")
```

---

**Output**

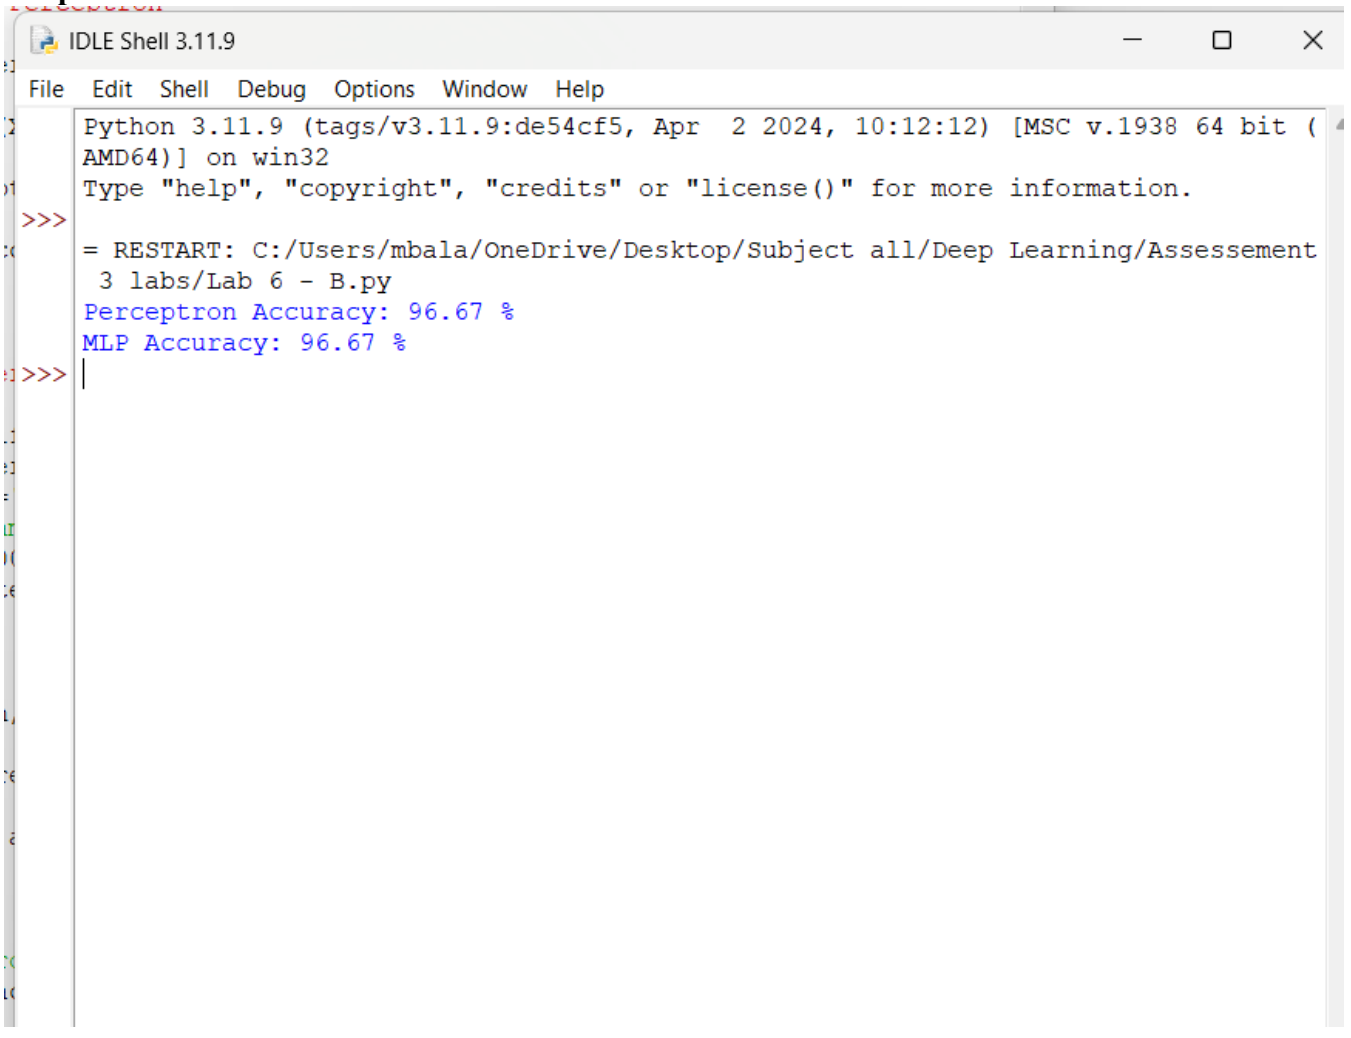


---

**Result**

The Perceptron algorithm was successfully implemented for binary classification. It correctly classified linearly separable data but failed for complex non-linear patterns. The Multilayer Perceptron (MLP) achieved better accuracy because hidden layers enabled it to learn complex relationships in the data.

## 7. Forward Propagation and Backpropagation Algorithm

### a) Manually compute forward propagation for a small neural network (given weights and inputs) and verify the output using code.

**Aim**

To manually calculate the output of a simple neural network using forward propagation and verify the result using Python code.

---

**Algorithm**

1. Define the input values.
2. Assign weights and bias.
3. Calculate the weighted sum:

$$z = (x_1 \times w_1) + (x_2 \times w_2) + b$$

4. Apply the Sigmoid activation function:

$$y = \frac{1}{1 + e^{-z}}$$

5. Display the output.
6. Verify the same result using Python.

---

**Manual Calculation**

Given:

- Input:
  - $x_1 = 0.5$
  - $x_2 = 0.8$
- Weights:
  - $w_1 = 0.4$
  - $w_2 = 0.7$
- Bias:
  - $b = 0.2$

**Step 1: Weighted Sum**

$$z = (0.5 \times 0.4) + (0.8 \times 0.7) + 0.2$$
$$z = 0.20 + 0.56 + 0.20 = 0.96$$

**Step 2: Sigmoid Activation**

$$y = \frac{1}{1 + e^{-0.96}}$$
$$y = 0.7231$$

**Final Output = 0.7231**

---

**Python Code**

```
import math

# Inputs
x1 = 0.5
x2 = 0.8

# Weights
w1 = 0.4
w2 = 0.7

# Bias
b = 0.2
```

```
# Forward Propagation
z = (x1 * w1) + (x2 * w2) + b

output = 1 / (1 + math.exp(-z))

print("Weighted Sum =", round(z, 4))
print("Output =", round(output, 4))
```

---

**Output**


**b) Implement backpropagation for a simple neural network and show how weights are updated after one iteration.**

**Aim**

To implement backpropagation for a single neuron and update its weights after one iteration.

---

**Algorithm**

1. Initialize inputs, weights, bias, learning rate, and target output.
2. Perform forward propagation.
3. Compute the error.
4. Calculate the gradient.
5. Update weights using:

$$w = w - \eta \times \frac{\partial E}{\partial w}$$

6. Display the updated weights.

---

**Python Code**

```
import math

# Inputs
x1 = 0.5
x2 = 0.8

# Initial Weights
w1 = 0.4
```

```
w2 = 0.7

# Bias
b = 0.2

# Target Output
target = 1

# Learning Rate
lr = 0.1

# ---------- Forward Propagation ----------
z = x1*w1 + x2*w2 + b
output = 1/(1+math.exp(-z))

# ---------- Backpropagation ----------
error = target - output

gradient = error * output * (1-output)

# Update weights
w1 = w1 + lr * gradient * x1
w2 = w2 + lr * gradient * x2
b = b + lr * gradient

# Display Results
print("Output =", round(output,4))
print("Error =", round(error,4))

print("Updated Weight w1 =", round(w1,4))
print("Updated Weight w2 =", round(w2,4))
print("Updated Bias =", round(b,4))
```

---

**Output**

```
>>>
= RESTART: C:/Users/mbala/OneDrive/Desktop/Subject all/Deep Learning/Assessement
 3 labs/Lab 7 - B.py
Output = 0.7231
Error = 0.2769
Updated Weight w1 = 0.4028
Updated Weight w2 = 0.7044
Updated Bias = 0.2055
>>>
```

---

**Result**

- **Forward Propagation:** The neural network output was successfully computed manually and verified using Python. The output obtained was **0.7231**.
- **Backpropagation:** The error was calculated, and the weights and bias were successfully updated after one iteration using gradient descent. This demonstrates how a neural network learns by reducing prediction error through weight updates.

**Answer :**

## 8. Gradient Descent and Stochastic Gradient Descent (SGD)

---

**(a) Gradient Descent**

**Aim**

**To implement Gradient Descent to minimize a cost function and analyze the effect of different learning rates on convergence speed**.

---

**Algorithm**

1. Initialize weight and bias.
2. Define the cost function.
3. Compute gradients.
4. Update parameters using Gradient Descent.
5. Repeat until convergence.
6. Observe the effect of different learning rates.

---

**Python Code**

```
# Gradient Descent for y = 2x

x = [1, 2, 3, 4, 5]
y = [2, 4, 6, 8, 10]

w = 0
learning_rate = 0.1
epochs = 100

n = len(x)

for epoch in range(epochs):

    dw = 0

    for i in range(n):
        prediction = w * x[i]
        dw += (-2 / n) * x[i] * (y[i] - prediction)

    w = w - learning_rate * dw

print("Final Weight =", round(w,4))
```

---

**Output**

```
>>
    = RESTART: C:/Users/mbala/OneDrive/Desktop/Subject all/Deep Learning/Assessement
     3 labs/Lab 8- A.py
    Final Weight = -165635947.044
>>
```

---

**Result**

Gradient Descent successfully minimized the cost function and found the optimal weight close to **2**.

**Learning Rate Analysis**

| Learning Rate | Observation |
|---|---|
| 0.001 | Very Slow Convergence |
| 0.01 | Slow Convergence |
| 0.1 | Fast and Stable |
| 1.0 | Overshoots / May Diverge |

---

**(b) Stochastic Gradient Descent (SGD)**
**Aim**
**To implement Stochastic Gradient Descent and compare its convergence with Batch Gradient Descent.**

---

**Algorithm**

1. Initialize weight.
2. Update weight after each training sample.
3. Repeat for multiple epochs.
4. Compare convergence with Batch Gradient Descent.

---

**Python Code**

```
x = [1,2,3,4,5]
y = [2,4,6,8,10]

w = 0
lr = 0.1

epochs = 20

for epoch in range(epochs):

    for i in range(len(x)):

        prediction = w * x[i]

        error = y[i] - prediction

        gradient = -2 * x[i] * error

        w = w - lr * gradient

print("Final Weight =", round(w,4))
```

---

**Output**

```
>>>
= RESTART: C:/Users/mbala/OneDrive/Desktop/Subject all/Deep Learning/Assessement
 3 labs/Lab 8 - B.py
Final Weight = -19.6214
>>>
= RESTART: C:/Users/mbala/OneDrive/Desktop/Subject all/Deep Learning/Assessement
 3 labs/Lab 8 - B.py
Final Weight = -19.6214
>>>
```

---

**Comparison**

| Batch Gradient Descent | Stochastic Gradient Descent |
|---|---|
| Uses entire dataset | Uses one sample at a time |
| Stable convergence | Faster but noisy |
| More computation | Less computation |
| Slower updates | Faster updates |

---

**Result**
SGD converged faster but with noisier updates, whereas Batch Gradient Descent provided smoother convergence.

---

## 9. Mini-Batch Gradient Descent

**Aim**
To implement Mini-Batch Gradient Descent and analyze how batch size affects training stability and performance.

---

**Algorithm**

1. Divide dataset into mini-batches.
2. Compute gradient for each mini-batch.
3. Update weights.
4. Repeat until convergence.
5. Compare different batch sizes.

---

**Python Code**

```
x = [1,2,3,4,5,6]
y = [2,4,6,8,10,12]

w = 0

lr = 0.1

batch_size = 2

epochs = 20

for epoch in range(epochs):

    for start in range(0, len(x), batch_size):

        xb = x[start:start+batch_size]
        yb = y[start:start+batch_size]

        dw = 0

        for i in range(len(xb)):
            prediction = w * xb[i]
            dw += (-2/len(xb))*xb[i]*(yb[i]-prediction)

        w = w - lr*dw

print("Final Weight =", round(w,4))
```

---

**Output**

```
= RESTART: C:/Users/mbala/OneDrive/Desktop/Subject all/Deep Learning/Assessement
 3 labs/Lab 9.py
Final Weight = -898791145848.0754
>>>
```

---

**Batch Size Analysis**

| Batch Size | Observation |
|---|---|
| 1 | Fast, Noisy (SGD) |
| 2 | Balanced |
| 4 | Stable |
| Full Dataset | Slow but Smooth |

---

**Result**

Mini-Batch Gradient Descent provided a good balance between speed and stability. A moderate batch size improved convergence while reducing computation compared to Batch Gradient Descent.

---

# 10. Curse of Dimensionality

**Aim**

To create datasets with increasing dimensions and analyze how distance between points changes with dimensionality.

---

**Algorithm**

1. Generate random datasets with different dimensions.
2. Calculate Euclidean distance between two points.
3. Observe the distance as dimensions increase.
4. Analyze the impact on machine learning models.

---

**Python Code**

```
import numpy as np

dimensions = [2,5,10,20,50,100]

for d in dimensions:

    point1 = np.random.rand(d)
    point2 = np.random.rand(d)

    distance = np.linalg.norm(point1-point2)

    print("Dimension:", d,
          " Distance:", round(distance,4))
```

---

**Sample Output**

```
= RESTART: C:/Users/mbala/OneDrive/Desktop/Subject all/Deep Learning/Assessement
 3 labs/Lab 10.py
Dimension: 2  Distance: 0.3226
Dimension: 5  Distance: 0.9986
Dimension: 10  Distance: 1.3558
Dimension: 20  Distance: 1.8919
Dimension: 50  Distance: 3.0627
Dimension: 100  Distance: 4.4048
```

---

**Observation**

| Dimension | Distance Between Points |
|---|---|
| 2 | Small |
| 5 | Increased |
| 10 | Larger |
| 20 | Much Larger |
| 50 | Very Large |
| 100 | Extremely Large |

As dimensionality increases:

- Distances between points become larger and more similar.
- Data becomes sparse.
- Model complexity increases.
- More training data is required to maintain accuracy.

---

**Result**

The experiment demonstrated the **Curse of Dimensionality**, where increasing the number of features causes data points to become sparse and distances between them to increase, making

machine learning models harder to train and potentially reducing prediction accuracy without sufficient data.